



INTERVIEW PREP • 2026

30

Gen AI Interview Questions **You Must Know**

If you can answer all 30, you're ready. Grouped into the 6 areas every interviewer digs into.



Foundations



LLMs



Prompting



Fine-Tuning



RAG



Production



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 1 · Q1-5

1

How is Generative AI different from discriminative AI?

2

Explain the Transformer architecture — why did it replace RNNs?

3

What is self-attention, and why is it $O(n^2)$ in sequence length?

4

Diffusion vs GANs vs VAEs — when would you use each?

5

What are embeddings, and what makes two of them "similar"?



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 2 • Q6-11

6

What actually happens in one forward pass of an LLM?

7

Encoder-only vs decoder-only vs encoder-decoder — an example of each?

8

What are scaling laws, and what did Chinchilla change?

9

Why do LLMs hallucinate, and how do you reduce it?

10

Explain temperature, top-k, and top-p sampling.

11

What is the context window, and why is it a hard limit?



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 3 · Q12-16

12

Zero-shot vs few-shot vs chain-of-thought prompting.

13

What is in-context learning — is the model really "learning"?

14

How do you design a system prompt for a production app?

15

What is prompt injection, and how do you defend against it?

16

How do you get reliable structured (JSON) output from an LLM?



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 4 • Q17-21

17

Prompt engineering vs RAG vs fine-tuning — how do you choose?

18

What is LoRA, and why is it cheaper than full fine-tuning?

19

Explain RLHF end-to-end.

20

What is catastrophic forgetting, and how do you avoid it?

21

What data do you actually need to fine-tune well?



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 5 • Q22-26

22

Walk through a RAG pipeline from document to answer.

23

What is chunking, and why does chunk size matter?

24

Vector vs keyword vs hybrid search — when each?

25

What is re-ranking, and where does it fit in the pipeline?

26

How do you evaluate a RAG system's quality?



Naved Khan

Senior Gen-AI Engineer

+ Follow



SECTION 6 • Q27-30

27

How do you evaluate an output when there's no single right answer?

28

What is "LLM-as-a-judge", and what are its pitfalls?

29

How do you cut LLM inference cost and latency in production?

30

What guardrails do you add around an LLM before shipping?



Bonus: For each, be ready to name a real tradeoff — that's what separates senior from junior answers.



Naved Khan

Senior Gen-AI Engineer

+ Follow

Answers that **get offers**

Knowing the definition isn't enough. Do these four things:



Lead with the tradeoff. Every choice costs something — name it before they ask.



Use a concrete example. "In a support-bot RAG, I'd chunk by section because..."



Connect it to production. Talk cost, latency, evals — not just theory.



Say what you don't know. "I'd verify X" beats confidently guessing.



Naved Khan

Senior Gen-AI Engineer

+ Follow

Save this for your prep 



Naved Khan

Senior Gen-AI Engineer



Like



Comment



Share



Save

Follow for more Gen-AI interview prep →