

GEN AI ENGINEER • 2026

75+ GenAI Interview Questions

they'll actually ask you

compiled from **534+ AI job postings** & real
2026 interview loops

Swipe to know how many you can answer.
Most freeze after 50 →

ROUND1: FUNDAMENTALS (CORE CONCEPTS)

1. What is a Large Language Model?
2. What is tokenization?
3. What is an embedding?
4. What is RAG?
5. What is a vector database?
6. What is prompt engineering?
7. What is the context window?
8. What is hallucination?
9. Fine-tuning vs RAG vs Prompt Engineering
10. Encoder-only vs Decoder-only vs Encoder-Decoder
11. Semantic caching
12. Temperature vs Top-p vs Top-k
13. Quantization
14. Latency vs Throughput
15. Speculative Decoding
16. Zero-shot vs One-shot vs Few-shot
17. Embedding Model vs Generation Model
18. Streaming vs Batch Inference

ROUND3: AGENTS, FRAMEWORKS & ADVANCED GENAI

- 37. LoRA vs QLoRA
- 38. ReAct Pattern
- 39. LLM-as-Judge
- 40. MCP
- 41. vLLM
- 42. Prevent Infinite Agent Loops
- 43. Agent Evaluation
- 44. Context Engineering
- 45. DSPy
- 46. Structured Outputs
- 47. Attention Sink
- 48. Prompt Versioning
- 49. LangChain vs LangGraph
- 50. Needle-in-a-Haystack Test
- 51. RLHF vs DPO
- 52. Chain-of-Thought
- 53. System Prompts & Security
- 54. Constitutional AI
- 55. Debugging Agents
- 56. OpenAI Agents SDK vs LangGraph vs CrewAI
- 57. Chain-of-Verification
- 58. Agent Memory Staleness
- 59. LLM Gateway vs Orchestration Framework
- 60. Graceful Degradatio

ROUND4: PRODUCTION SYSTEMS, SECURITY & SYSTEM DESIGN

61. OWASP LLM Top 10
62. Low-Latency LLM Systems
63. API Failures & Rate Limits
64. User Satisfaction Drop Investigation
65. On-Prem Deployment
66. Duplicate Email Agent Bug
67. Reducing LLM Costs
68. Exponential Backoff Coding
69. Build RAG Pipeline
70. Tool Calling & Validation
71. Performance Improvement Story
72. Quality vs Cost Tradeoff
73. Production Incident
74. Staying Current
75. Open-Weight Models vs Frontier Models
76. PII Handling
77. A/B Testing Prompt Changes
78. Model Distillation
79. Handling Traffic Spikes
80. GenAI Product Metrics Dashboard

ROUND2: RAG & RETRIEVAL ENGINEERING

19. RAG failure modes
20. Hybrid Search
21. RAGAS Metrics
22. Multi-hop RAG architecture
23. Lost-in-the-middle problem
24. Multi-tenant RAG
25. Handling ambiguous queries
26. Legal Document Q&A System
27. HNSW vs IVF
28. Prompt Compression
29. Knowledge Graphs & GraphRAG
30. Metadata Filtering
31. Reranker vs Retriever
32. Retrieval-based Few-shot Prompting
33. Document Q&A vs Conversational RAG
34. Contradicting Documents in RAG
35. Preventing stale information in RAG
36. pgvector vs Dedicated Vector DB